

SCTP for Gophers

Why You Should Try SCTP In Your Next Golang Back-end Project

Apr 2019

Alexey Naidyonov
CSA, ITooLabs

What is SCTP?

A transport layer IP protocol (like TCP or UDP), used mostly to deliver telephone network signalling over IP.

SCTP is designed to transport Public Switched Telephone Network (PSTN) signaling messages over IP networks, but is capable of broader applications.

[RFC 4960 "Stream Control Transmission Protocol"](https://tools.ietf.org/html/rfc4960) (<https://tools.ietf.org/html/rfc4960>)

- Message-oriented
- Reliable
- Multi-stream
- Multi-home

Disclaimer

This talk does not cover:

- SCTP over UDP
- WebRTC Data Channels
- usrsctp
- ... or anything alike

SCTP Support

- Linux/FreeBSD/Solaris
- No (native) Windows support (usrctp)
- No (native) Darwin support (usrctp, kext)
- Weak NAT traversal support on consumer devices

Suitable for server side, not so for client applications

SCTP Features (1/2)

- Message-oriented
- Message boundary preservation
- Reliable data transfer w/SACK
- Unordered data delivery option
- Multi-stream support
- Message bundling
- Path MTU discovery, message fragmentation

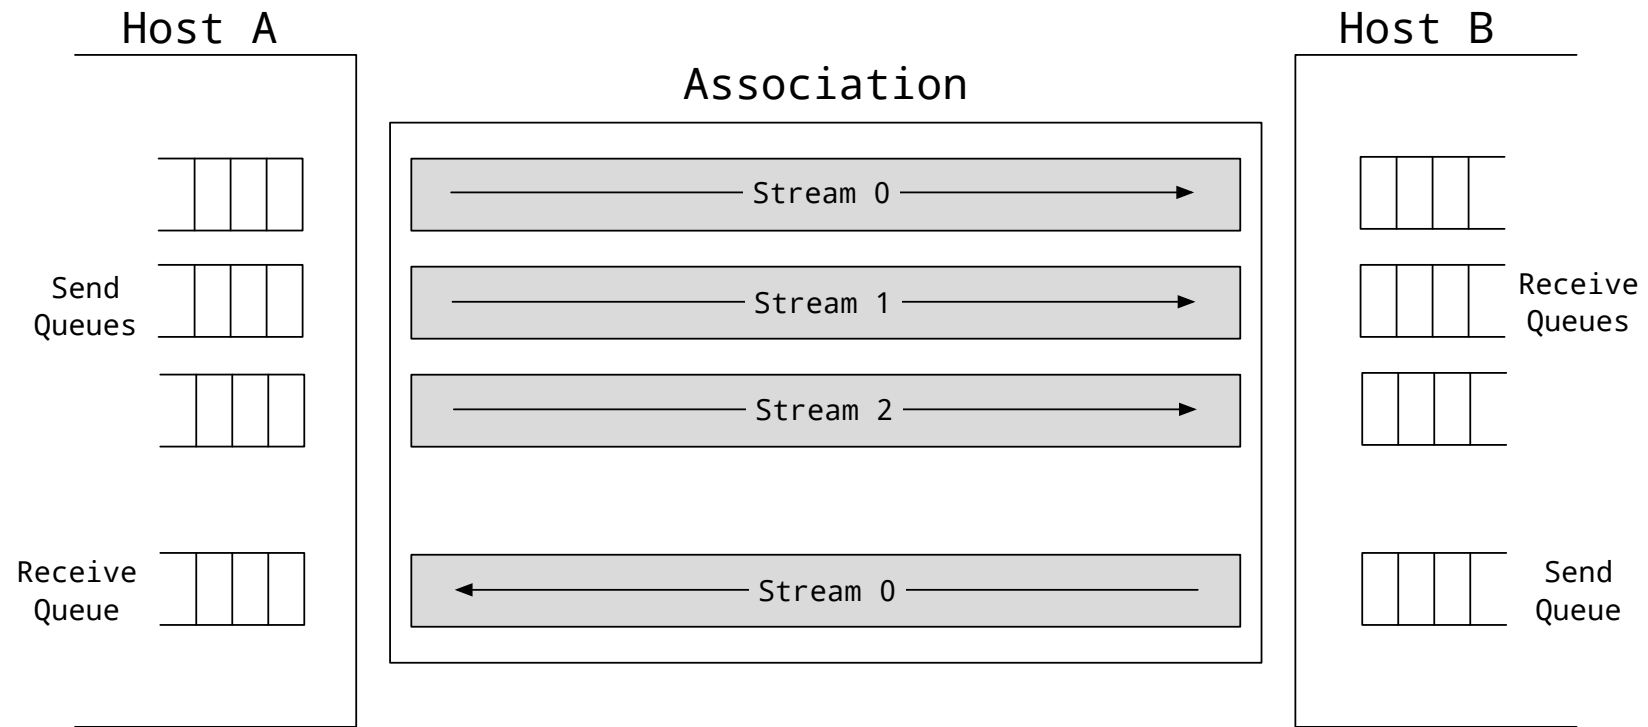
SCTP Features (2/2)

- Congestion control and avoidance
- Multi-homing support
- Built-in heartbeat (reachability check)
- 4-way handshake with security cookies to prevent connection flood attack
- 3-way shutdown
- One-to-one and one-to-many connections
- Extensible

SCTP Model

- **Endpoint:** an addressable logical endpoint, represented by single port number and one or more IP addresses
- **Association:** a logical relationship between two endpoints (\approx connection)
- **Stream:** a logical unidirectional channel transporting applications messages.
- **Message:** a user message delivered over stream
- **Packet:** PDU comprises of header and multiple **chunks**
- **Chunk:** a basic data unit containing either protocol control information (control chunk) or user data (data chunk)
- **Transmission Sequence Number (TSN):** 32-bit sequence number attached to every data chunk
- **Stream Sequence Number (SSN):** 16-bit sequence number attached to every message of a stream,


SCTP Multi-stream support



SCTP Multi-stream support

- Every stream is independent (no head-of-line blocking)
- Messages within a stream are ordered using the SSN
- Message fragments have the same SSN
- Message might be unordered
- Number of streams are negotiated at INIT
- Number of outbound streams: how many streams we plan to send?
- Max number of inbound streams: how many streams we wish to accept?
- Number of streams might be re-negotiated

SCTP Multi-homing support





SCTP Multi-homing support

- Multi-homing means multiple transport addresses
- Makes sense when different IP addresses are bound to different network paths
- One path is *primary*, in case of failure SCTP falls back to *alternate* path
- API allows to send data to a specific transport address
- Heartbeat checks performed over all paths
- Addresses might be added or removed on the fly ([RFC5061](https://tools.ietf.org/html/rfc5061) (https://tools.ietf.org/html/rfc5061))

SCTP Association init and shutdown

- 4-way handshake



- 3-way shutdown
- No half-closed state

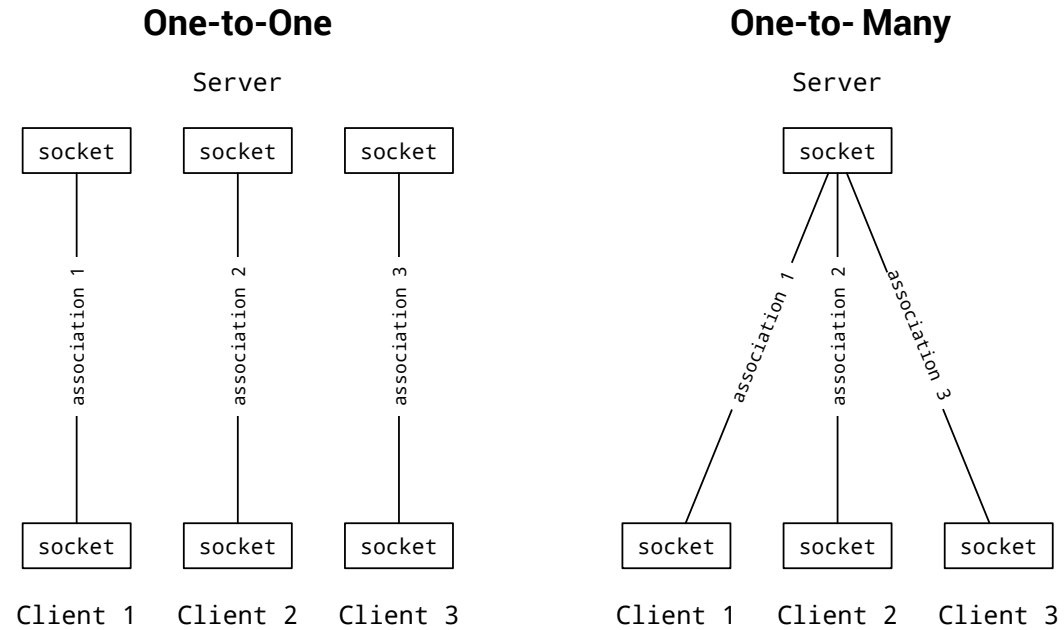
SCTP Flow control

- Guarantees receiver buffer will never overflow
- Acknowledges successful receipt of DATA chunks (using SACK chunks)
- Reports Gaps and Duplicated TSNs
- Guarantees fast recovery

SCTP API

- Not trivial due to model complexity
- Defined by [RFC 6458 "Socket API Extensions for SCTP"](https://tools.ietf.org/html/rfc6458) (<https://tools.ietf.org/html/rfc6458>)
- One-to-one socket
- One-to-many socket
- Event notifications

SCTP API



- sctp_peeloff API branches a single one-to-many association to one-to-one socket
- SCTP_SENDALL flag forces broadcast to all one-to-many associations

SCTP Go API

- Multiple cgo
- Native github.com/nperez-messagebird/sctp (<https://github.com/nperez-messagebird/sctp>)
- Native ITooLabs, not open (yet)

SCTP Go API problems

- Go native networking is way too native
- Go `syscall.RawConn` (starting go 1.9)

```
type RawConn interface {  
    // Control invokes f on the underlying connection's file  
    // descriptor or handle.  
    Control(f func(fd uintptr)) error  
  
    // Read invokes f on the underlying connection's file  
    // descriptor or handle; f is expected to try to read from the  
    // file descriptor.  
    Read(f func(fd uintptr) (done bool)) error  
  
    // Write is like Read but for writing.  
    Write(f func(fd uintptr) (done bool)) error  
}
```

- (still not enough)

SCTP Go API: the minimal set of runtime calls required

```
//go:linkname sysSocket net.sysSocket
func sysSocket(family, sotype, proto int) (int, error)

//go:linkname newNetFD net.newFD
func newNetFD(sysfd, family, sotype int, net string) (unsafe.Pointer, error)

//go:linkname initNetFD net.(*netFD).init
func initNetFD(netfd unsafe.Pointer) error

//go:linkname setDeadline net.(*netFD).SetDeadline
func setDeadline(netfd unsafe.Pointer, t time.Time) error

//go:linkname seReadDeadline net.(*netFD).SetReadDeadline
func setReadDeadline(netfd unsafe.Pointer, t time.Time) error

//go:linkname setWriteDeadline net.(*netFD).SetWriteDeadline
func setWriteDeadline(netfd unsafe.Pointer, t time.Time) error

//go:linkname setsockopt syscall.setsockopt
func setsockopt(s int, level int, name int, val unsafe.Pointer, vallen uintptr) (err error)

//go:linkname getsockopt syscall.getsockopt
func getsockopt(s int, level int, name int, val unsafe.Pointer, vallen *_Socklen) (err error)
```

SCTP Go API:

Waiting for [#15021](https://github.com/golang/go/issues/15021) "Allow registrations for new socket type API"

 **golang / go**

Watch 3,282

Star 56,238

Fork 7,669

<> Code

 Issues 4,373

 Pull requests 110

 Wiki

 Insights

syscall: allow registration of new socket types for package net #15021

 **Closed** mdlayher opened this issue on 29 Mar 2016 · 37 comments



mdlayher commented on 29 Mar 2016

Member + 👤 ...

Overview

At this time, there is no mechanism for socket types outside of the standard library to access the runtime network poller. This proposal, if accepted, would enable a resolution to issue [#10565](#). This would enable packages outside of the standard library to take advantage of the runtime network poller, instead of implementing their own network polling mechanism.

Assignees

No one assigned

Labels

Proposal

Proposal-Accepted

early-in-cycle

SCTP Go API (github.com/nperez-messagebird/sctp)

```
// Addr represents SCTP end point address
//
// Single SCTP end point might be bound to multiple IP addresses.
//
type SCTPAddr struct {
    IPAddrs    []net.IPAddr
    Port       int
}

// Resolves SCTP Address. Multiple IP addresses are separated by "/"
// (i.e. 10.50.1.1/10.50.2.1:2904, 10.50.1.1/[2001:db8::1]:2904)
//
func ResolveSCTPAddr(addressFamily SCTPAddressFamily, addrs string) (*SCTPAddr, error)
```

SCTP Go API (github.com/nperez-messagebird/sctp)

```
type SCTPSocketMode int

const (
    OneToOne = SCTPSocketMode(iota)
    OneToMany
)

type InitMsg struct {
    NumOstreams    uint16
    MaxInstreams   uint16
    MaxAttempts     uint16
    MaxInitTimeout uint16
}

func NewSCTPListener(laddr *SCTPAddr, init InitMsg, mode SCTPSocketMode) (*SCTPListener, error)
func NewSCTPConnection(laddr, raddr *SCTPAddr, options InitMsg, mode SCTPSocketMode) (*SCTPConn, error)
```

SCTP Go API (github.com/nperez-messagebird/sctp)

```
type SndRcvInfo struct {
    Stream  uint16 // stream id
    SSN     uint16 // stream sequence number
    Flags   uint16 // flags (i.e. SCTP_UNORDERDED)
    _       uint16
    PPID    uint32 // supplied by peer app
    Context uint32 // message context (supplied by app)
    TTL     uint32 // time to live (ms)
    TSN     uint32 // TSN
    CumTSN  uint32 // Cumulative TSN
    AssocID int32  // association id
}

func SCTPRead(b []byte) (int, *OOMessage, int, error)
func (o *OOMessage) GetSndRcvInfo() *SndRcvInfo

func SCTPWrite(b []byte, info *SndRcvInfo) (int, error)
```

SCTP Go API (github.com/nperez-messagebird/sctp)

```
// read (one-to-one)
func (ln *SCTPConn) SCTPRead(b []byte) (n int, oob *OOBMessage, flags int, err error)

// write (one-to-one)
func (ln *SCTPConn) SCTPWrite(b []byte, info *SndRcvInfo) (n int, err error)
```

SCTP Go API (github.com/nperez-messagebird/sctp)

```
// Accept (one-to-one mode)
func (ln *SCTPListener) AcceptSCTP() (*SCTPConn, error)

// read (one-to-many mode)
func (ln *SCTPListener) SCTPRead(b []byte) (n int, oob *OOBMessage, flags int, err error)

// write (one-to-many mode)
func (ln *SCTPListener) SCTPWrite(b []byte, info *SndRcvInfo) (n int, err error)
```


SCTP Go API (github.com/nperez-messagebird/sctp)

```
func SubscribeEvents(int evts) error

const (
    SCTP_EVENT_DATA_IO = 1 << iota
    SCTP_EVENT_ASSOCIATION
    SCTP_EVENT_ADDRESS
    SCTP_EVENT_SEND_FAILURE
    SCTP_EVENT_PEER_ERROR
    SCTP_EVENT_SHUTDOWN
    SCTP_EVENT_PARTIAL_DELIVERY
    SCTP_EVENT_ADAPTATION_LAYER
    SCTP_EVENT_AUTHENTICATION
    SCTP_EVENT_SENDER_DRY
)
```

SCTP Go API (IToolLabs)

```
func NewSCTPSocket(network string, mode SocketMode) (*Socket, error)

func NewSCTPListener(network string, laddr string, mode SocketMode) (*Listener, error)

// Conn is an association, not the socket itself
func DialSCTP(network string, laddr *Addr, mode SocketMode) (*Conn, error)

func (l *Listener) Accept() (*Conn, error)

func (c *Conn) PeelOff() error

// Stream returns a stream handler
func (c *Conn) IStream(sid uint16) (*IStream, error) {...}
func (c *Conn) OStream(sid uint16) (*OStream, error) {...}

func (s *OStream) WriteMsg([]byte) error
func (s *IStream) RecvMsg([]byte) error
```

SCTP Go API (Events)


```
// Represents SCTP notification event
type Event interface {
    Type() uint16 // SCTP event type
    Flags() uint16 // SCTP event flags
    Len() uint32 // SCTP event length, including header
}

// EventPeerAddrChange
// EventAssocChange
// AuthenticationEvent
// ShutdownEvent
// SenderDryEvent
// PartialDeliveryEvent
// SendFailureEvent
// ...

func (s *Socket) SubscribeEvents(events EventKind...)
func (s *Socket) Events() <-chan Event
```

Multi-homing demo

- Environment: VirtualBox + Docker + OpenVSWitch
- 2 virtual networks (10.10.100/24, 10.10.200/24)
- 2 containers (client, server)



SCTP: Where to use

- A distributed multi-node multi-service application (seem to be ~80% of all Go applications out there)
- Strict requirements for high availability and high resiliency
- Controlled multi-path networking environment

Thank you

Apr 2019

Tags: SCTP, networking (#ZgotmplZ)

Alexey Naidyonov

CSA, ITooLabs

anaidyonov@itoolabs.com (mailto:anaidyonov@itoolabs.com)

<https://itoolabs.com/> (https://itoolabs.com/)

[@growler](http://twitter.com/growler) (http://twitter.com/growler)